

In [1]: *# Install necessary packages*

```
!pip install pandas  
!pip install numpy  
!pip install scipy  
!pip install scikit-learn  
!pip install matplotlib  
!pip install seaborn
```

Requirement already satisfied: pandas in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (2.2.1)

Requirement already satisfied: numpy<2,>=1.23.2 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from pandas) (1.26.4)

Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from pandas) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from pandas) (2023.3.post1)

Requirement already satisfied: tzdata>=2022.7 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from pandas) (2024.1)

Requirement already satisfied: six>=1.5 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)

Requirement already satisfied: numpy in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (1.26.4)

Requirement already satisfied: scipy in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (1.12.0)

Requirement already satisfied: numpy<1.29.0,>=1.22.4 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from scipy) (1.26.4)

Requirement already satisfied: scikit-learn in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (1.4.1.post1)

Requirement already satisfied: numpy<2.0,>=1.19.5 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from scikit-learn) (1.26.4)

Requirement already satisfied: scipy>=1.6.0 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from scikit-learn) (1.12.0)

Requirement already satisfied: joblib>=1.2.0 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from scikit-learn) (1.3.2)

Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from scikit-learn) (3.3.0)

Requirement already satisfied: matplotlib in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (3.8.3)

Requirement already satisfied: contourpy>=1.0.1 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib) (1.2.0)

Requirement already satisfied: cyclor>=0.10 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib) (4.49.0)

Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib) (1.4.5)

Requirement already satisfied: numpy<2,>=1.21 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib) (1.26.4)

Requirement already satisfied: packaging>=20.0 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib) (23.1)

Requirement already satisfied: pillow>=8 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib) (10.2.0)

Requirement already satisfied: pyparsing>=2.3.1 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib) (3.1.2)

Requirement already satisfied: python-dateutil>=2.7 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib) (2.8.2)

Requirement already satisfied: six>=1.5 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from python-dateutil>=2.7->matplotlib) (1.16.0)

Requirement already satisfied: seaborn in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (0.13.2)

Requirement already satisfied: numpy!=1.24.0,>=1.20 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from seaborn) (1.26.4)

Requirement already satisfied: pandas>=1.2 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from seaborn) (2.2.1)

Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from seaborn) (3.8.3)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.2.0)
Requirement already satisfied: cycler>=0.10 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (4.49.0)
Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.4.5)
Requirement already satisfied: packaging>=20.0 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (23.1)
Requirement already satisfied: pillow>=8 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (10.2.0)
Requirement already satisfied: pyparsing>=2.3.1 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from pandas>=1.2->seaborn) (2023.3.post1)
Requirement already satisfied: tzdata>=2022.7 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from pandas>=1.2->seaborn) (2024.1)
Requirement already satisfied: six>=1.5 in c:\users\josue\anaconda3\envs\d206performanceassessment\lib\site-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.4->seaborn) (1.16.0)

```
In [2]: # Standard imports
import numpy as np
import pandas as pd
from sklearn.preprocessing import scale
from sklearn.decomposition import PCA

import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [3]: # Load data set into Pandas dataframe
churn_df = pd.read_csv('C:\\Users\\josue\\Desktop\\WGU\\D206\\churn_raw_data.csv')
```

```
In [4]: #display churn dataframe to know it actually loaded up.
#A. Research Question how many people over the age of 62 have an active service and
churn_df
```

Out[4]:

	Unnamed: 0	CaseOrder	Customer_id		Interaction	City
0	1	1	K409198		aa90260b-4141-4a24-8e36-b04ce1f4f77b	Point Baker
1	2	2	S120509	fb76459f-c047-4a9d-8af9-e0f7d4ac2524		West Branch
2	3	3	K191035		344d114c-3736-4be5-98f7-c72c281e2d35	Yamhill
3	4	4	D90850		abfa2b40-2d43-4994-b15a-989b8c79e311	Del Mar
4	5	5	K662701		68a861fd-0d20-4e51-a587-8a90407ee574	Needville
...
9995	9996	9996	M324793		45deb5a2-ae04-4518-bf0b-c82db8dbe4a4	Mount Holly
9996	9997	9997	D861732		6e96b921-0c09-4993-bbda-a1ac6411061a	Clarksville
9997	9998	9998	I243405	e8307ddf-9a01-4fff-bc59-4742e03fd24f		Mobeetie
9998	9999	9999	I641617	3775ccfc-0052-4107-81ae-9657f81ecd3		Carrollton
9999	10000	10000	T38070		9de5fb6e-bd33-4995-aec8-f01d0172a499	Clarkesville

10000 rows × 52 columns

```
In [5]: #List of dataframe columns
df = churn_df.columns
print(df)
```

```
Index(['Unnamed: 0', 'CaseOrder', 'Customer_id', 'Interaction', 'City',
      'State', 'County', 'Zip', 'Lat', 'Lng', 'Population', 'Area',
      'Timezone', 'Job', 'Children', 'Age', 'Education', 'Employment',
      'Income', 'Marital', 'Gender', 'Churn', 'Outage_sec_perweek', 'Email',
      'Contacts', 'Yearly_equip_failure', 'Techie', 'Contract', 'Port_modem',
      'Tablet', 'InternetService', 'Phone', 'Multiple', 'OnlineSecurity',
      'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
      'StreamingMovies', 'PaperlessBilling', 'PaymentMethod', 'Tenure',
      'MonthlyCharge', 'Bandwidth_GB_Year', 'item1', 'item2', 'item3',
      'item4', 'item5', 'item6', 'item7', 'item8'],
      dtype='object')
```

```
In [6]: #remove unnamed column and display first five records
df = churn_df.drop(churn_df.columns[0], axis = 1)
df.head()
```

Out[6]:	CaseOrder	Customer_id	Interaction	City	State	County	Zip	
0	1	K409198	aa90260b-4141-4a24-8e36-b04ce1f4f77b	Point Baker	AK	Prince of Wales-Hyder	99927	56
1	2	S120509	fb76459f-c047-4a9d-8af9-e0f7d4ac2524	West Branch	MI	Ogemaw	48661	46
2	3	K191035	344d114c-3736-4be5-98f7-c72c281e2d35	Yamhill	OR	Yamhill	97148	45
3	4	D90850	abfa2b40-2d43-4994-b15a-989b8c79e311	Del Mar	CA	San Diego	92014	32
4	5	K662701	68a861fd-0d20-4e51-a587-8a90407ee574	Needville	TX	Fort Bend	77461	25

5 rows × 51 columns

```
In [7]: #rename last 8 survey columns to more descriptive names rather than item
df.rename(columns = { 'item1':'Responses',
                      'item2':'Solutions',
                      'item3':'Replacements',
                      'item4':'Reliability',
                      'item5':'Options',
                      'item6':'Respectfulness',
                      'item7':'Courteous',
                      'item8':'Listening'},
          inplace=True)
```

```
In [8]: #amount of records and columns of the dataset
df.shape
```

Out[8]: (10000, 51)

```
In [9]: #describe Churn dataset columns and rows
df.describe()
```

```
Out[9]:
```

	CaseOrder	Zip	Lat	Lng	Population	Children
count	10000.00000	10000.000000	10000.000000	10000.000000	10000.000000	7505.000000
mean	5000.50000	49153.319600	38.757567	-90.782536	9756.562400	2.095936
std	2886.89568	27532.196108	5.437389	15.156142	14432.698671	2.154758
min	1.00000	601.000000	17.966120	-171.688150	0.000000	0.000000
25%	2500.75000	26292.500000	35.341828	-97.082812	738.000000	0.000000
50%	5000.50000	48869.500000	39.395800	-87.918800	2910.500000	1.000000
75%	7500.25000	71866.500000	42.106908	-80.088745	13168.000000	3.000000
max	10000.00000	99929.000000	70.640660	-65.667850	111850.000000	10.000000

8 rows × 7 columns

```
In [10]: #remove variables from the dataset that do not help answer the question then display
df_stats = df.drop(columns=['CaseOrder', 'Zip', 'Lat', 'Lng'])
df_stats.describe()
```

```
Out[10]:
```

	Population	Children	Age	Income	Outage_sec_perweek	Churn
count	10000.000000	7505.000000	7525.000000	7510.000000	10000.000000	10000.000000
mean	9756.562400	2.095936	53.275748	39936.762226	11.452955	0.2575
std	14432.698671	2.154758	20.753928	28358.469482	7.025921	0.4371
min	0.000000	0.000000	18.000000	740.660000	-1.348571	0.000000
25%	738.000000	0.000000	35.000000	19285.522500	8.054362	0.000000
50%	2910.500000	1.000000	53.000000	33186.785000	10.202896	0.000000
75%	13168.000000	3.000000	71.000000	53472.395000	12.487644	0.000000
max	111850.000000	10.000000	89.000000	258900.700000	47.049280	1.000000

```
In [11]: #calculate churn rate
df.Churn.value_counts() / len(df)
```

```
Out[11]: Churn
No      0.735
Yes     0.265
Name: count, dtype: float64
```

```
In [12]: #review data types ( numerical => "int64", "float64", while categorical data is =>
df.dtypes
```

```
Out[12]: CaseOrder          int64
Customer_id        object
Interaction         object
City               object
State              object
County             object
Zip                int64
Lat                float64
Lng                float64
Population          int64
Area               object
Timezone           object
Job                object
Children           float64
Age                float64
Education           object
Employment          object
Income             float64
Marital            object
Gender             object
Churn              object
Outage_sec_perweek float64
Email              int64
Contacts            int64
Yearly_equip_failure int64
Techie             object
Contract           object
Port_modem         object
Tablet             object
InternetService    object
Phone              object
Multiple           object
OnlineSecurity      object
OnlineBackup        object
DeviceProtection    object
TechSupport         object
StreamingTV         object
StreamingMovies     object
PaperlessBilling    object
PaymentMethod       object
Tenure              float64
MonthlyCharge       float64
Bandwidth_GB_Year   float64
Responses           int64
Solutions           int64
Replacements        int64
Reliability         int64
Options             int64
Respectfulness      int64
Courteous           int64
Listening           int64
dtype: object
```

```
In [13]: # Re-validate column data types and missing values
df.columns.to_series().groupby(df.dtypes).groups
```

```
Out[13]: {int64: ['CaseOrder', 'Zip', 'Population', 'Email', 'Contacts', 'Yearly equip_failure', 'Responses', 'Solutions', 'Replacements', 'Reliability', 'Options', 'Respectfulness', 'Courteous', 'Listening'], float64: ['Lat', 'Lng', 'Children', 'Age', 'Income', 'Outage_sec_perweek', 'Tenure', 'MonthlyCharge', 'Bandwidth_GB_Year'], object: ['Customer_id', 'Interaction', 'City', 'State', 'County', 'Area', 'Timezone', 'Job', 'Education', 'Employment', 'Marital', 'Gender', 'Churn', 'Techie', 'Contract', 'Port_modem', 'Tablet', 'InternetService', 'Phone', 'Multiple', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies', 'PaperlessBilling', 'PaymentMethod']}
```

```
In [14]: #find missing values
df.isnull()
```

```
Out[14]:
```

	CaseOrder	Customer_id	Interaction	City	State	County	Zip	Lat	Lng	Popula
0	False	False	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	False	False	
...	
9995	False	False	False	False	False	False	False	False	False	
9996	False	False	False	False	False	False	False	False	False	
9997	False	False	False	False	False	False	False	False	False	
9998	False	False	False	False	False	False	False	False	False	
9999	False	False	False	False	False	False	False	False	False	

10000 rows × 51 columns

```
In [15]: #access only rows from dataframe containing missing values
df.isnull().any(axis=1)
```

```
Out[15]: 0      True
1      False
2      True
3      False
4      False
...
9995   True
9996   True
9997   True
9998   False
9999   True
Length: 10000, dtype: bool
```

```
In [16]: #display specific columns with no value i.e NA
df.isna().any()
```



```
Out[16]: CaseOrder      False
         Customer_id    False
         Interaction     False
         City            False
         State           False
         County          False
         Zip             False
         Lat             False
         Lng             False
         Population      False
         Area            False
         Timezone        False
         Job             False
         Children        True
         Age             True
         Education       False
         Employment      False
         Income          True
         Marital         False
         Gender          False
         Churn           False
         Outage_sec_perweek False
         Email           False
         Contacts        False
         Yearly_equip_failure False
         Techie          True
         Contract        False
         Port_modem      False
         Tablet          False
         InternetService  True
         Phone           True
         Multiple        False
         OnlineSecurity  False
         OnlineBackup    False
         DeviceProtection False
         TechSupport     True
         StreamingTV     False
         StreamingMovies  False
         PaperlessBilling False
         PaymentMethod    False
         Tenure          True
         MonthlyCharge    False
         Bandwidth_GB_Year True
         Responses       False
         Solutions       False
         Replacements     False
         Reliability      False
         Options          False
         Respectfulness  False
         Courteous        False
         Listening         False
         dtype: bool
```

```
In [17]: # store null values in a variable to get total counts
```

```
d_nulls = df.isnull().sum()
```

```
print(d_nulls)
```

CaseOrder	0
Customer_id	0
Interaction	0
City	0
State	0
County	0
Zip	0
Lat	0
Lng	0
Population	0
Area	0
Timezone	0
Job	0
Children	2495
Age	2475
Education	0
Employment	0
Income	2490
Marital	0
Gender	0
Churn	0
Outage_sec_perweek	0
Email	0
Contacts	0
Yearly equip_failure	0
Techie	2477
Contract	0
Port_modem	0
Tablet	0
InternetService	2129
Phone	1026
Multiple	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	991
StreamingTV	0
StreamingMovies	0
PaperlessBilling	0
PaymentMethod	0
Tenure	931
MonthlyCharge	0
Bandwidth_GB_Year	1021
Responses	0
Solutions	0
Replacements	0
Reliability	0
Options	0
Respectfulness	0
Courteous	0
Listening	0

dtype: int64

```
In [18]: # Store rows with missing values in a new variable
rows_with_missing_values = df.isnull().any(axis=1)
df[rows_with_missing_values]
```

```
Out[18]:
```

	CaseOrder	Customer_id	Interaction	City	State	Count
0	1	K409198	aa90260b-4141-4a24-8e36-b04ce1f4f77b	Point Baker	AK	Prince c Wales-Hyde
2	3	K191035	344d114c-3736-4be5-98f7-c72c281e2d35	Yamhill	OR	Yamhi
5	6	W303516	2b451d12-6c2b-4cea-a295-ba1d6bced078	Fort Valley	GA	Peac
6	7	U335188	6630d501-838c-4be4-a59c-6f58c814ed6a	Pioneer	TN	Sco
7	8	V538685	70ddaa89-b726-49dc-9022-2d655e4c7936	Oklahoma City	OK	Oklahom
...	
9994	9995	P175475	c60df12b-a50b-4397-ae57-98381a0d3960	West Kill	NY	Green
9995	9996	M324793	45deb5a2-ae04-4518-bf0b-c82db8dbe4a4	Mount Holly	VT	Rutlan
9996	9997	D861732	6e96b921-0c09-4993-bbda-a1ac6411061a	Clarksville	TN	Montgomer
9997	9998	I243405	e8307ddf-9a01-4fff-bc59-4742e03fd24f	Mobeetie	TX	Wheelc
9999	10000	T38070	9de5fb6e-bd33-4995-aec8-f01d0172a499	Clarkesville	GA	Habershar

8316 rows × 51 columns

```
In [19]: #Examine columns for miss spelled words in categorical variables using unquie() met
df['Employment'].unique()
```

```
Out[19]: array(['Part Time', 'Retired', 'Student', 'Full Time', 'Unemployed'],
dtype=object)
```

```
In [20]: df['Area'].unique()
```

```
Out[20]: array(['Urban', 'Suburban', 'Rural'], dtype=object)
```

```
In [21]: df['Timezone'].unique()
```

```
Out[21]: array(['America/Sitka', 'America/Detroit', 'America/Los_Angeles',
               'America/Chicago', 'America/New_York', 'America/Puerto_Rico',
               'America/Denver', 'America/Menominee', 'America/Phoenix',
               'America/Indiana/Indianapolis', 'America/Boise',
               'America/Kentucky/Louisville', 'Pacific/Honolulu',
               'America/Indiana/Petersburg', 'America/Nome', 'America/Anchorage',
               'America/Indiana/Knox', 'America/Juneau', 'America/Toronto',
               'America/Indiana/Winamac', 'America/Indiana/Vincennes',
               'America/North_Dakota/New_Salem', 'America/Indiana/Tell_City',
               'America/Indiana/Marengo', 'America/Ojinaga'], dtype=object)
```

```
In [22]: #get total count of how many unquie items there are
len(df['Job'].unique())
```

```
Out[22]: 639
```

```
In [23]: df['Children'].unique()
```

```
Out[23]: array([nan, 1., 4., 0., 3., 2., 7., 5., 9., 6., 10., 8.])
```

```
In [24]: #find the unique amount of ages then get a count for how many differnet age's ther
df['Age'].unique()
```

```
Out[24]: array([68., 27., 50., 48., 83., nan, 49., 86., 23., 56., 30., 39., 63.,
               60., 61., 52., 75., 77., 47., 70., 69., 45., 40., 82., 26., 25.,
               66., 72., 41., 44., 43., 84., 59., 31., 51., 58., 73., 33., 42.,
               81., 87., 54., 67., 46., 24., 20., 71., 32., 29., 80., 53., 79.,
               65., 35., 34., 74., 55., 76., 57., 38., 78., 19., 36., 88., 62.,
               37., 28., 22., 85., 89., 18., 21., 64.])
```

```
In [25]: age_range = df['Age'].unique()
print(sorted(age_range))
```

```
[23.0, 25.0, 26.0, 27.0, 30.0, 31.0, 39.0, 40.0, 41.0, 43.0, 44.0, 45.0, 47.0, 48.0,
49.0, 50.0, 51.0, 52.0, 59.0, 61.0, 68.0, 83.0, nan, 18.0, 19.0, 20.0, 21.0, 22.0, 2
4.0, 28.0, 29.0, 32.0, 33.0, 34.0, 35.0, 36.0, 37.0, 38.0, 42.0, 46.0, 53.0, 54.0, 5
5.0, 56.0, 57.0, 58.0, 60.0, 62.0, 63.0, 64.0, 65.0, 66.0, 67.0, 69.0, 70.0, 71.0, 7
2.0, 73.0, 74.0, 75.0, 76.0, 77.0, 78.0, 79.0, 80.0, 81.0, 82.0, 84.0, 85.0, 86.0, 8
7.0, 88.0, 89.0]
```

```
In [26]: #finding more unquie values in other columns
df['Education'].unique()
```

```
Out[26]: array(["Master's Degree", 'Regular High School Diploma',
               'Doctorate Degree', 'No Schooling Completed', "Associate's Degree",
               "Bachelor's Degree", 'Some College, Less than 1 Year',
               'GED or Alternative Credential',
               'Some College, 1 or More Years, No Degree',
               '9th Grade to 12th Grade, No Diploma',
               'Nursery School to 8th Grade', 'Professional School Degree'],
              dtype=object)
```

```
In [27]: df['Employment'].unique()
```

```
Out[27]: array(['Part Time', 'Retired', 'Student', 'Full Time', 'Unemployed'],
              dtype=object)
```

```
In [28]: df['Marital'].unique()
```

```
Out[28]: array(['Widowed', 'Married', 'Separated', 'Never Married', 'Divorced'],
              dtype=object)
```

```
In [29]: df['Gender'].unique()
```

```
Out[29]: array(['Male', 'Female', 'Prefer not to answer'], dtype=object)
```

```
In [30]: df['Contract'].unique()
```

```
Out[30]: array(['One year', 'Month-to-month', 'Two Year'], dtype=object)
```

```
In [31]: df['PaymentMethod'].unique()
```

```
Out[31]: array(['Credit Card (automatic)', 'Bank Transfer(automatic)',
              'Mailed Check', 'Electronic Check'], dtype=object)
```

```
In [32]: # Display any duplicate rows in the dataframe.
         data_duplicates = df.loc[df.duplicated()]
         print(data_duplicates)
```

Empty DataFrame

Columns: [CaseOrder, Customer_id, Interaction, City, State, County, Zip, Lat, Lng, Population, Area, Timezone, Job, Children, Age, Education, Employment, Income, Marital, Gender, Churn, Outage_sec_perweek, Email, Contacts, Yearly equip_failure, Techie, Contract, Port_modem, Tablet, InternetService, Phone, Multiple, OnlineSecurity, OnlineBackup, DeviceProtection, TechSupport, StreamingTV, StreamingMovies, PaperlessBilling, PaymentMethod, Tenure, MonthlyCharge, Bandwidth_GB_Year, Responses, Solutions, Replacements, Reliability, Options, Respectfulness, Courteous, Listening]

Index: []

[0 rows x 51 columns]

```
In [33]: data_nulls = df_stats.isnull().sum()
         print(data_nulls)
```

Customer_id	0
Interaction	0
City	0
State	0
County	0
Population	0
Area	0
Timezone	0
Job	0
Children	2495
Age	2475
Education	0
Employment	0
Income	2490
Marital	0
Gender	0
Churn	0
Outage_sec_perweek	0
Email	0
Contacts	0
Yearly_equip_failure	0
Techie	2477
Contract	0
Port_modem	0
Tablet	0
InternetService	2129
Phone	1026
Multiple	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	991
StreamingTV	0
StreamingMovies	0
PaperlessBilling	0
PaymentMethod	0
Tenure	931
MonthlyCharge	0
Bandwidth_GB_Year	1021
Responses	0
Solutions	0
Replacements	0
Reliability	0
Options	0
Respectfulness	0
Courteous	0
Listening	0

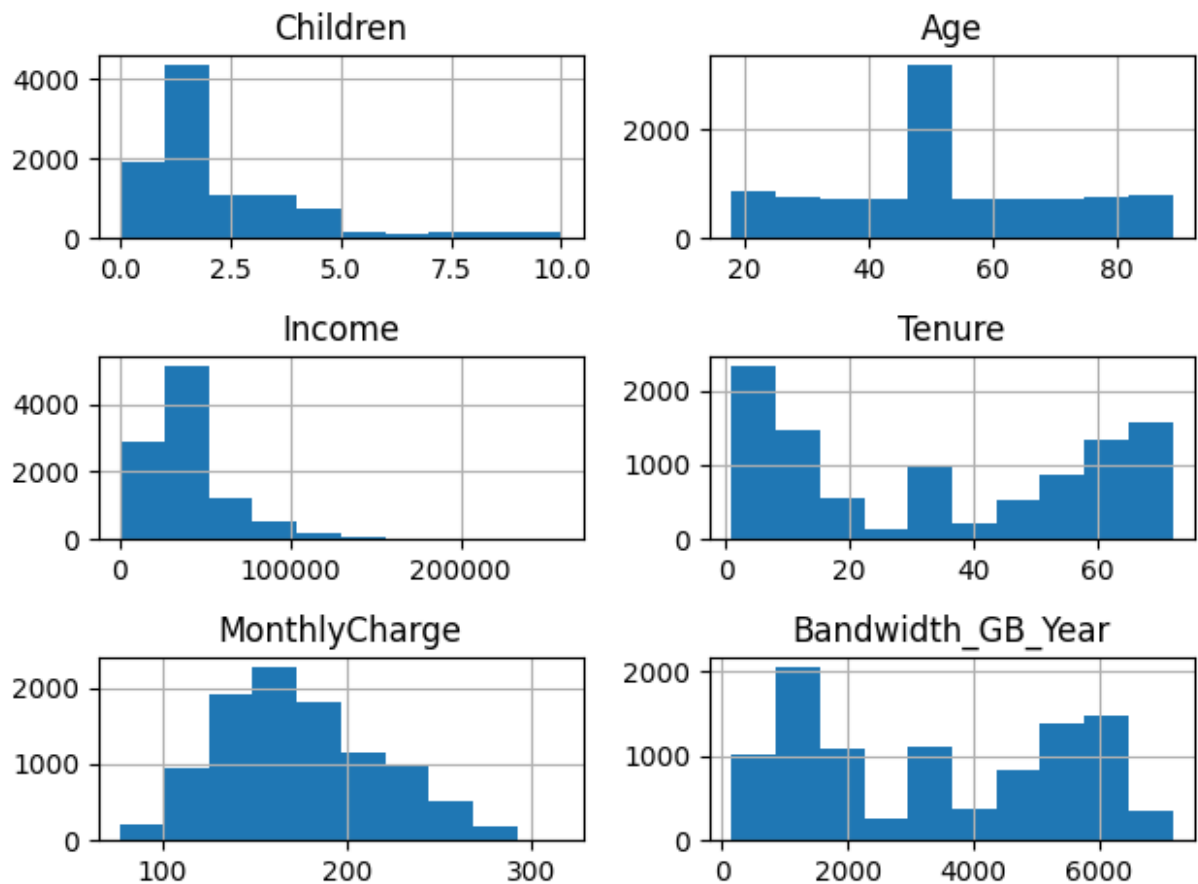
dtype: int64

```
In [34]: #impute missing fields for variables with median or mean
df_stats['Children'] = df['Children'].fillna(df['Children'].median())
df_stats['Age'] = df['Age'].fillna(df['Age'].median())
df_stats['Income'] = df['Income'].fillna(df['Income'].median())
df_stats['Tenure'] = df['Tenure'].fillna(df['Tenure'].median())
df_stats['Bandwidth_GB_Year'] = df['Bandwidth_GB_Year'].fillna(df['Bandwidth_GB_Yea
```

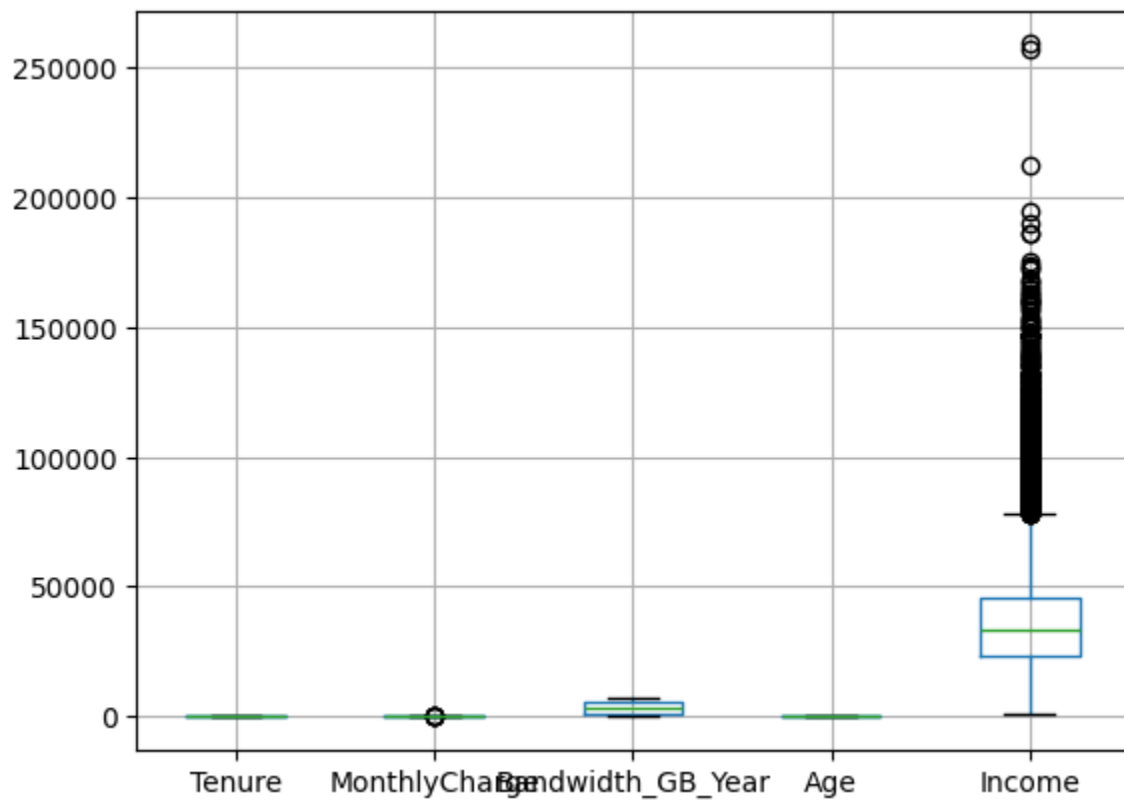
```
In [35]: data_nulls = df_stats.isnull().sum()
print(data_nulls)
```

```
Customer_id      0
Interaction       0
City              0
State             0
County            0
Population        0
Area              0
Timezone          0
Job               0
Children          0
Age               0
Education         0
Employment        0
Income            0
Marital           0
Gender            0
Churn             0
Outage_sec_perweek 0
Email             0
Contacts          0
Yearly_equip_failure 0
Techie            2477
Contract         0
Port_modem        0
Tablet            0
InternetService   2129
Phone             1026
Multiple          0
OnlineSecurity    0
OnlineBackup      0
DeviceProtection  0
TechSupport       991
StreamingTV       0
StreamingMovies   0
PaperlessBilling  0
PaymentMethod     0
Tenure            0
MonthlyCharge     0
Bandwidth_GB_Year 0
Responses         0
Solutions         0
Replacements      0
Reliability       0
Options           0
Respectfulness    0
Courteous         0
Listening         0
dtype: int64
```

```
In [36]: # Create histograms of important variables
df_stats[['Children', 'Age', 'Income', 'Tenure', 'MonthlyCharge', 'Bandwidth_GB_Yea
plt.savefig('churn_pyplot.jpg')
plt.tight_layout()
```




```
In [77]: # Create a boxplot of user duration, payment, useage,income & age variables
df_stats.boxplot(['Tenure', 'MonthlyCharge', 'Bandwidth_GB_Year', 'Age', 'Income'])
plt.savefig('churn_boxplots.jpg')
```





```
In [38]: #isolate monlthy charge where we can take a better look at the outliers
df_stats.boxplot(['MonthlyCharge'])
```

Out[38]: <Axes: >



```
In [55]: # Lets take a closer look at income to see the outliers here
df_stats.boxplot(['Tenure', 'Income'])
```

Out[55]: <Axes: >



```
In [56]: #extract Clean dataset
df_stats.to_csv('churn_cleaned.csv')
```

```
In [59]: #reload cleaned data & remove all variable except user services payment info and su
churn_user = pd.read_csv('churn_cleaned.csv')
```

```
In [60]: #remove all but the last 11 service related variables for the PCA
data = churn_user.loc[:, 'Tenure':'Listening']
data.head()
```

```
Out[60]:
```

	Tenure	MonthlyCharge	Bandwidth_GB_Year	Responses	Solutions	Replacements	Re
0	6.795513	171.449762	904.536110	5	5	5	
1	1.156681	242.948015	800.982766	3	4	3	
2	15.754144	159.440398	2054.706961	4	4	2	
3	17.087227	120.249493	2164.579412	4	4	4	
4	1.670972	150.761216	271.493436	4	4	4	

```
In [61]: #Import Scikit Learn PCA application
from sklearn.decomposition import PCA
```

```
In [62]: #Normalize the data
churn_normalized = (data - data.mean()) / data.std()
```

```
In [63]: #Select number of components to extract
pca = PCA(n_components = data.shape[1])
```

```
In [65]: #Create a List of PCA names
churn_numeric = data[['Tenure', 'MonthlyCharge', 'Bandwidth_GB_Year', 'Responses',
                     'Solutions', 'Replacements', 'Reliability', 'Options',
                     'Respectfulness', 'Courteous', 'Listening']]


pcs_names = []
for i, col in enumerate(churn_numeric.columns):
    pcs_names.append('PC' + str(i + 1))
print(pcs_names)

['PC1', 'PC2', 'PC3', 'PC4', 'PC5', 'PC6', 'PC7', 'PC8', 'PC9', 'PC10', 'PC11']
```

```
In [66]: #Call PCA application & convert the dataset of 11 variables into a dataset of 11 co
pca.fit(churn_normalized)
churn_pca = pd.DataFrame(pca.transform(churn_normalized),
                        columns = pcs_names)
```


```
In [67]: #For a scree plot import matplotlib & seaborn libraries
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [69]: #Run the scree plot
plt.plot(pca.explained_variance_ratio_)
plt.xlabel('Number of Components')
plt.ylabel(' Variance')
plt.show();
```



```
In [70]: #Extract the eigenvalues
cov_matrix = np.dot(churn_normalized.T, churn_normalized) / data.shape[0]
eigenvalues = [np.dot(eigenvector.T, np.dot(cov_matrix, eigenvector)) for eigenvector in eigenvectors]
```

```
In [71]: #Plot the eigenvalues
plt.plot(eigenvalues)
plt.xlabel('Number of Components')
plt.ylabel('Eigenvalue')
plt.show();
```



```
In [72]: for pc, var in zip(pcs_names, np.cumsum(pca.explained_variance_ratio_)):
          print(pc, var)
```

```
PC1 0.2679266062396318
PC2 0.44053732880229535
PC3 0.589144452224861
PC4 0.6801579842425111
PC5 0.751321705437251
PC6 0.814311709410132
PC7 0.8681970531672304
PC8 0.9171387362588809
PC9 0.9610122820002168
PC10 0.9905736363144829
PC11 0.9999999999999998
```

```
In [73]: # we see that 86% of variance is explained by 7 components
# Create a rotation
rotation = pd.DataFrame(pca.components_.T, columns = pcs_names, index = churn_numbers)
print(rotation)
```

	PC1	PC2	PC3	PC4	PC5	PC6	\
Tenure	-0.010403	0.701838	-0.072209	-0.063594	0.005683	-0.011155	
MonthlyCharge	0.000317	0.041147	-0.014151	0.996995	-0.022136	0.015231	
Bandwidth_GB_Year	-0.012166	0.703079	-0.074222	0.004399	0.009590	0.003466	
Responses	0.458932	0.031325	0.281154	0.018568	-0.070233	-0.119149	
Solutions	0.434134	0.042559	0.282404	0.007508	-0.106632	-0.169752	
Replacements	0.400639	0.034665	0.281118	-0.019631	-0.173742	-0.255336	
Reliability	0.145799	-0.050367	-0.567815	-0.010310	-0.171334	-0.483328	
Options	-0.175633	0.066334	0.587335	-0.000047	0.135949	0.060124	
Respectfulness	0.405207	-0.012680	-0.183447	0.004596	-0.062342	0.064609	
Courteous	0.358342	-0.003886	-0.181697	-0.027959	-0.182406	0.806166	
Listening	0.308925	-0.017396	-0.131173	0.015574	0.931612	-0.011133	

	PC7	PC8	PC9	PC10	PC11
Tenure	0.007419	-0.011527	0.006935	0.003286	-0.705445
MonthlyCharge	-0.018038	-0.004316	0.023690	-0.013785	-0.047865
Bandwidth_GB_Year	0.003701	-0.002364	-0.008068	0.008529	0.706925
Responses	-0.045963	0.025431	-0.240574	0.793237	-0.004306
Solutions	-0.065414	0.074400	-0.592131	-0.573832	-0.002217
Replacements	-0.146887	-0.396333	0.673088	-0.177665	0.014933
Reliability	-0.443353	0.431528	0.087207	0.018301	0.002283
Options	-0.209767	0.693861	0.265474	-0.042012	-0.002514
Respectfulness	0.757954	0.402835	0.230319	-0.063972	0.001604
Courteous	-0.379136	0.067889	0.067293	-0.040946	-0.006875
Listening	-0.113297	-0.045132	0.046107	-0.042251	-0.002357

```
In [74]: # Output Loadings for components
loadings = pd.DataFrame(pca.components_.T,
                        columns = pcs_names,
                        index = data.columns)

loadings
```

Out[74]:

	PC1	PC2	PC3	PC4	PC5	PC6	F
Tenure	-0.010403	0.701838	-0.072209	-0.063594	0.005683	-0.011155	0.0074
MonthlyCharge	0.000317	0.041147	-0.014151	0.996995	-0.022136	0.015231	-0.0180
Bandwidth_GB_Year	-0.012166	0.703079	-0.074222	0.004399	0.009590	0.003466	0.0037
Responses	0.458932	0.031325	0.281154	0.018568	-0.070233	-0.119149	-0.0459
Solutions	0.434134	0.042559	0.282404	0.007508	-0.106632	-0.169752	-0.0654
Replacements	0.400639	0.034665	0.281118	-0.019631	-0.173742	-0.255336	-0.1468
Reliability	0.145799	-0.050367	-0.567815	-0.010310	-0.171334	-0.483328	-0.4433
Options	-0.175633	0.066334	0.587335	-0.000047	0.135949	0.060124	-0.2097
Respectfulness	0.405207	-0.012680	-0.183447	0.004596	-0.062342	0.064609	0.7579
Courteous	0.358342	-0.003886	-0.181697	-0.027959	-0.182406	0.806166	-0.3791
Listening	0.308925	-0.017396	-0.131173	0.015574	0.931612	-0.011133	-0.1132

```
In [75]: #reduced dataset & print 3 components
```

```
churn_reduced = churn_pca.iloc[ : , 0:3]
print(churn_reduced)
```

	PC1	PC2	PC3
0	1.923875	-1.421955	1.903125
1	-0.199798	-1.706801	0.538766
2	-0.667923	-0.985940	0.227390
3	0.046465	-0.730628	2.282040
4	1.326741	-1.924880	0.825729
...
9995	-2.097964	1.961837	0.104147
9996	1.917485	1.645946	0.611009
9997	1.431918	0.323573	0.028288
9998	2.011460	2.187756	-0.079864
9999	-2.266364	1.591986	-0.819973

[10000 rows x 3 columns]

In []: